



Nonlinear effects and generalized additive models

Roger Levy

30 April 2025

Inferring the trend



Generalized Additive Models (GAM) = other (e.g. linear regression)

Taking on nonlinearity: major ideas

- Use a **basis expansion**: a set of functions $\{h_i\}$ each of which differently transforms your original predictor X
- Learn the $Y \sim X$ as a linear comb. of the basis functions

$$f(x) = \sum_i \beta_i h_i(x)$$

$$y \sim \text{Exp}(f(x))$$




Some exponential-family function (one that you could use in a generalized linear model)


- Use **regularization** and **held-out accuracy evaluation** to guard against overfitting

Example basis expansions for *piecewise* fits

Piecewise Constant




Piecewise Linear




$$\begin{aligned} h_1 &= I(X < \xi_1) \\ h_2 &= I(\xi_1 < X \leq \xi_2) \\ h_3 &= I(\xi_2 < X \leq \xi_3) \\ h_4 &= h_1(X)X \\ h_5 &= h_2(X)X \\ h_6 &= h_3(X)X \end{aligned}$$

Continuous Piecewise Linear




Piecewise-linear Basis Function



$$\begin{aligned} h_1 &= 1 \\ h_2 &= X \\ h_3 &= (X - \xi_1)_+ \\ h_4 &= (X - \xi_2)_+ \end{aligned}$$

Example basis expansions for *piecewise* fits

Piecewise Cubic Polynomials



Popular
modeling
choice

Natural, or restricted, cubic splines

- Intuitively: a piecewise cubic function whose 0th, 1st, and second derivatives are continuous at all the **knots**
- Technically, for n knots:

$$h_1(X) = 1$$

$$h_2(X) = X$$


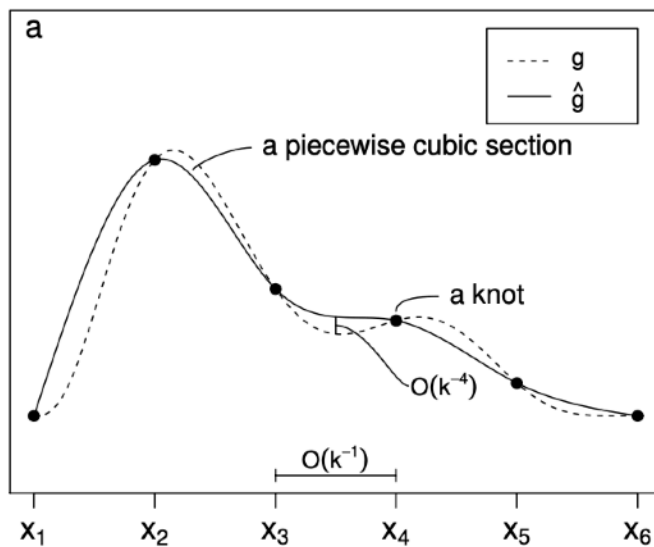

$$h_3(X) = X^2$$

$$h_4(X) = X^3$$


$$h_i(X) = (x - \xi_i)_+^3, \text{ for } i = 5, \dots, n + 4$$

- Powerful framework, but leaves open:
 - How many knots to use
 - Where to put the knots

Natural cubic splines, visualized



Example of fitted natural splines (5 d.f.)



From smooths to GAMs

- A **generalized additive model** (GAM) involving 2 predictors X and V for response Y looks like this:

$$y_i = \mathbf{A}_i\theta + f_1(x_i) + f_2(v_i) + f_3(x_i, v_i) + \epsilon_i$$

where each $f_i()$ is a **smooth function** of its arguments

- This is a highly general and expressive formalism, and cannot be fitted via maximum likelihood (why?)
- So in practice, GAMs are fitted by:
 - choosing a **penalization** or **regularization** component for the objective function;
 - choosing a set of **basis functions** by which to represent each smooth function;
 - Searching for a fit that optimizes the objective function

Penalization & cross-validation

- Penalized objective:

$$\|\mathbf{y} - \mathbf{X}\boldsymbol{\beta}\|^2 + \underbrace{\lambda_1 \boldsymbol{\beta}^\top \mathbf{S}_1 \boldsymbol{\beta}} + \underbrace{\lambda_2 \boldsymbol{\beta}^\top \mathbf{S}_2 \boldsymbol{\beta}}$$

Terms penalizing wiggleness for f_1 and f_2

How well data are predicted

- How to choose $\{\lambda_i\}$?
- Here: **cross-validation**

Strategy for getting proficient in GAMs

- Read Jones et al. 2013 Chapter 7, and do the exercises
- Read Wood 2017, starting with Chapter 4 and Chapter 7, and do the exercises
- To build expertise read Wood 2017 Chapter 6 and also Chapter 5
- Both these books are available online via MIT Libraries
- For the rest of this class, we will build a bit of expertise via worked examples
 - Data from Boyce & Levy, 2023
 - Material from the "GAMs in R" interactive course (<https://noamross.github.io/gams-in-r-course/>)